

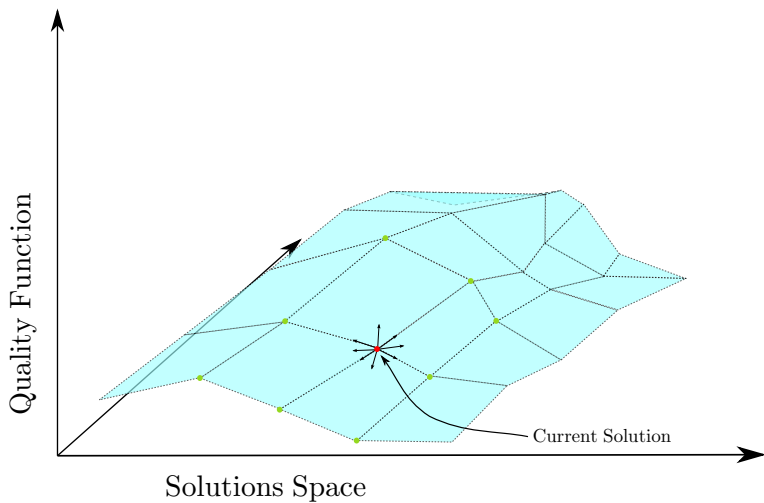
Local Search

- Sometimes we don't need the path that reaches a solution, we search in the space of solutions
- We want to obtain the best attainable solution in an affordable time (optimal is impossible)
- We have a function that evaluates the quality of the solution, this value is not related to a path cost
- Search is performed from an initial solution that we try to improve using actions
- The actions allow to move in the solution neighbourhood

Local search

- The heuristic function:
 - Approximates the quality of a solution (it is not a cost)
 - The goal is to optimize it (maximize or minimize)
 - Combines all the elements of the problem and its constraints (possibly using different weights for different elements)
 - There are no constraints about how the function can be, it only has to represent the quality relations among the solutions
 - It can be positive or negative

Local Search



Local Search

- The size of the space of solutions doesn't allow for an optimal solution search
- It is not possible to perform a systematic search
- The heuristic function is used to prune the space of solutions (solutions that don't need to be explored)
- Usually no history of the search path is stored (minimal space complexity)

Hill climbing

- First-choice Hill climbing
 - First action that improves the current solution is taken
- Steepest-ascent hill climbing, gradient search
 - The best action that improves the current solution is taken

Hill Climbing

Algorithm: Hill Climbing

Current \leftarrow initial state

End \leftarrow **false**

while not End do

 Successors \leftarrow generate_successor(Current)

 Successors \leftarrow sort_and_prune_bad_solutions(Successors, Current)

if *not empty?(Successors)* **then**

 Current \leftarrow best_successor(Successors)

else

 End \leftarrow **true**

end

end

- Only are considered successors those solutions with a heuristic function value better than the current solution (pruning of the space of solutions)
- A stack could be used to store the best successors to backtrack, but usually the space requirement are prohibitive
- The algorithm may not find any solution even when there are

Hill climbing


- The characteristics of the heuristic function and the initial solution determine the success and the time of the search
- The strategy of this algorithm may end the search in a solution that is only apparently the optimal
- Problems
 - Local optima: No neighbor solution has a better value
 - Plateaus: All neighbours have the same value
 - Ridge: A sequence of local optima

Hill climbing

- Possible solutions

- Backtrack to a previous solution and follow another path (it is only possible if we limit the memory used for backtracking, *Beam Search*)
- Restart the search from another initial solution looking for a better solution (*Random-restarting Hill-Climbing*)
- Use two or more actions to explore deeper the neighbourhood after making any decision (expensive in time and space)
- Parallel Hill-Climbing (for instance: divide the search space in regions and explore the most promising ones, possibly sharing information)


Hill Climbing - Example - Knapsack problem



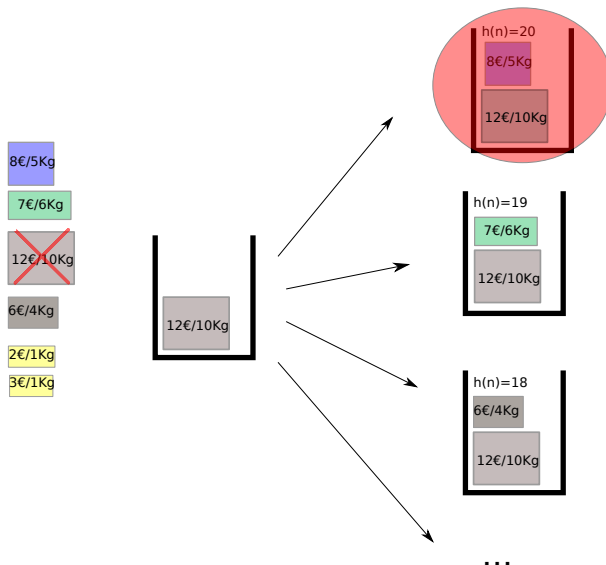
BY: Dake

- Solution: Any combination of objects inside the knapsack
- Initial solution: Empty knapsack
- Operators: Put objects in and take objects from the knapsack
- Heuristic Function: $\max \sum_i Value_i$ or $\max \sum_i \frac{Value_i}{Weight_i}$


Hill Climbing - Example - Knapsack problem



Hill Climbing - Example - Knapsack problem



Hill Climbing - Example - Knapsack problem




Other local search algorithms

- There are other local search algorithms with different inspirations like physics or biology:
 - **Simulated annealing:** Stochastic Hill-climbing inspired in the controlled cooling of metal alloys and substances dissolution
 - **Genetic Algorithms:** Parallel stochastic Hill-climbing inspired in the mechanism of natural selection
- But also Particle Swarm Optimization, Ant Colony Optimization, Intelligent Water Drop, Gravitational search algorithm, ...

Simulated Annealing

- Stochastic Hill-Climbing (a successor is randomly chosen from the neighbor solutions using a probability distribution, the successor could have worst evaluation than the current solution)
- A random walk of the space of solutions is performed
- Inspired in the physics of controlled annealing (crystallization, metal alloys tempering)
- A metal alloy or dissolution is heated at high temperatures and progressively cooled in a controlled way
- If the cooling process is adequate the minimal state of energy of the system is achieved (global minimum)

Simulated Annealing



BY NC ND DoITPoMS, University of Cambridge

Simulated Annealing - Methodology

- We have to identify the elements of the problem with the elements of the physics analogy
- **Temperature**, control parameter
- **Energy**, quality of the solution $f(n)$
- **Acceptance function**, allows to decide if to pick a successor solution
 - $\mathcal{F}(\Delta f, T)$, function of the temperature and the difference of quality between the current solution and the candidate solution
 - The lower temperature, the lower the chance to choose a successor with worst evaluation
- **Cooling strategy**, number of iterations to perform, how to lower the temperature and how many successors to explore each temperature step

Simulated annealing - canonical algorithm

Algorithm: Simulated Annealing

An initial temperature is chosen

while *temperature above zero* **do**

 // Random walk the space of solutions

for *the chosen number of iterations* **do**

 NewSol \leftarrow generate_random_successor(CurrentSol)

$\Delta E \leftarrow f(\text{CurrentSol}) - f(\text{NewSol})$

if $\Delta E > 0$ **then**

 CurrentSol \leftarrow NewSol

else

 with probability $e^{\Delta E/T}$: CurrentSol \leftarrow NewSol


end

end

 Reduce the temperature

end

Simulated Annealing



Simulated Annealing - Applications

- Used for combinatorial optimization problems (optimal configuration of a set of components) and continuous optimization (optimal in a N-dimensional space)
- Adequate for large sized problems in which the global optimal could be surrounded by lots of local optimums
- Adequate for problems where to find a discriminant heuristic is difficult (a random choice is as good as any other choice)
- Applications: TSP, Design of VLSI circuits
- Problems: To determine the value of the parameters of the algorithm requires experimentation (sometimes very extensive)


Simulated annealing - Example - TSP

- Traveler salesman problem (TSP): Search space $N!$
- Possible actions to change a solution: Inversions, translation, interchange
- An energy function (Sum of the distance among the cities, following the order in the solution)

$$E = \sum_{i=1}^n \sqrt{(x_i - x_{i+1})^2 + (y_i - y_{i+1})^2} + \sqrt{(x_N - x_1)^2 + (y_N - y_1)^2}$$

- We should define an initial temperature (experimentation)
- We should determine the number of iterations for each temperature step and how is the temperature decreased

Simulated annealing - Example - TSP



Genetic Algorithms

- Inspired in the mechanisms of natural selection
 - All living things adapt to environment because of the characteristics inherited from their parents
 - The probability of survival and reproduction is related to the quality of these characteristics (fitness of the individual to the environment)
 - The combination of good individuals could result in better adapted individuals
- We can translate this analogy to local search
 - The solutions are individuals
 - The fitness function indicates the quality of the solution
 - Combining good solutions we could obtain better solutions


Genetic algorithms (II)

To solve a problem using GA we need:

- To code the characteristics of the solutions (for example as a binary string)
- A function that measures the quality of a solution (fitness function)
- Operators that combine solutions to obtain new solutions (crossover operations)
- The number of individuals in the initial population
- An strategy about how to match the individuals

Genetic algorithms - Coding


- Usually the coding of the individuals is a binary string (it is not always the best)



- The coding defines the size of the search space and the crossover operators that are needed

Genetic algorithms - Operators

- The combination of individuals is done using crossover operators
- The basic operator is the one-point crossover
 - A cutting point in the coding is chosen randomly
 - The information of the two individuals is interchanged using this point



Genetic algorithms - Operators (II)

- There are other possibilities:
 - two-points crossover
 - random bit interchanging
 - specific operators depending on the coding
- Mutation operators:
 - Following the analogy to genetics and reproduction, sometimes a part of the gene changes randomly
 - The basic mutation operator is to change with a probability a randomly chosen bit in the coding

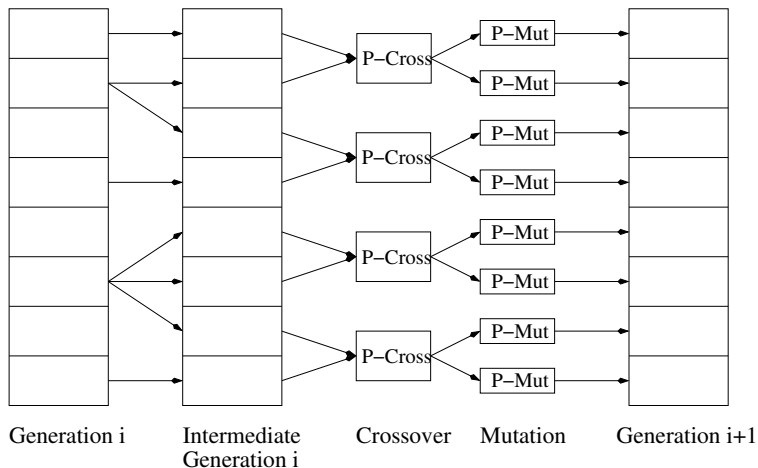
Genetic algorithms - Matching

- Each step in the search a new generation of individuals is obtained maintaining the size of the population constant (N)
- To obtain the next generation the individuals to combine (intermediate generation) are chosen following a criteria, for example:
 - Each individual is chosen with a probability proportional to its fitness
 - N random tournaments are performed among pairs of individuals, the individual that wins is chosen
 - A linear ranking among the individuals is defined using the fitness function
- Always some individuals will appear more than once and some will not appear at all

Genetic algorithms - canonical algorithm

- The basics steps of a GA are:
 - 1 N individuals are chosen from the current generation to form the intermediate generation (using a specific criteria)
 - 2 Individuals are paired and for each pair:
 - With probability ($P_{\text{crossover}}$) the crossover operator is applied and two new individuals are obtained
 - With probability (P_{mutation}) the new individuals are mutated
 - 3 These individuals conform the new generation
 - 4 Iterate until the population converges or a specific number of iterations is performed
- The crossover probability has a crucial influence in the diversity of the next generation
- The probability of mutation is always very low to avoid random search

Genetic algorithms - Canonical algorithm



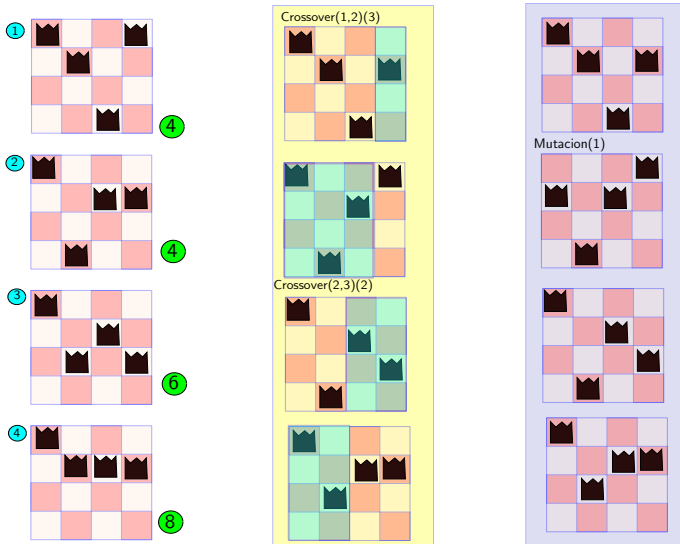
Genetic algorithms - Applications

- Are used virtually in any kind of problem
- They allow to solve problems that not have a known good heuristic
- Usually will perform worst than using hill climbing with a good heuristic
- Applications: Innumerable
- Problems: Coding the solutions, find the good parameters of the algorithm (size of the population, number of iterations, probability of crossover and mutation)
- In some problems GA perform poorly

Genetic algorithms - Example

- **N-queens problems**
- A solution can be coded as a binary string
- Individual= Concat($i=1 \dots N$; Binary(column(queen_{*i*})))
- Fitness function= number of pairs of queens that attack each other
- Crossover operator= one-point crossover
- Selection of the intermediate population: Proportional to the fitness function value
- Probability of crossover \rightarrow jexperiment!
- Probability of mutation \rightarrow jexperiment!
- Size of the initial population: ? (size of the search space n^n)

Genetic algorithms - Example - N queens



Genetic algorithms - Example - N queens

